



# Decision Trees

Department of Computer Languages and Systems

# Decision trees: introduction

## Objectives of a decision tree

- **inference, learning, tree generation:**
  - to achieve the best possible split, at leaf nodes, of the (training) samples of different classes
  - to build small trees
- **use, exploitation:**
  - to predict the class (classification)
  - to predict the value of a target variable (regression)

# Decision trees: introduction

Strategy to build a decision tree from a training set  $X$ :

- progressive splitting of the training set into smaller and smaller subsets
- recursive splitting of  $X$  based on the value of attributes
- creating a labelled leaf, when all instances of a subset belong to the same class
- pruning: a branch with a mixed subset of distinct classes is labelled with the majority class (the resulting subtree is pruned)

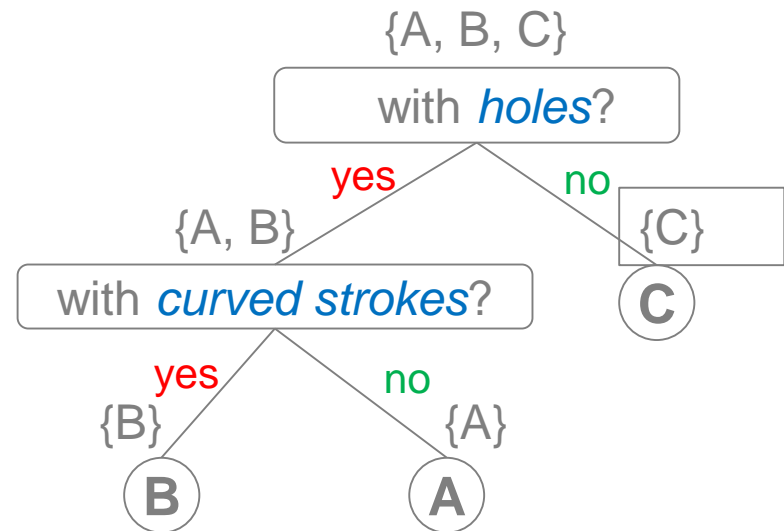
# Decision trees: an example

You want to define a **system to recognise 3 capital letters** (e.g., Arial) based on the presence or absence of 2 features: *holes* and *curved strokes*

class	holes	curved strokes
A	yes	no
B	yes	yes
C	no	yes

## comments:

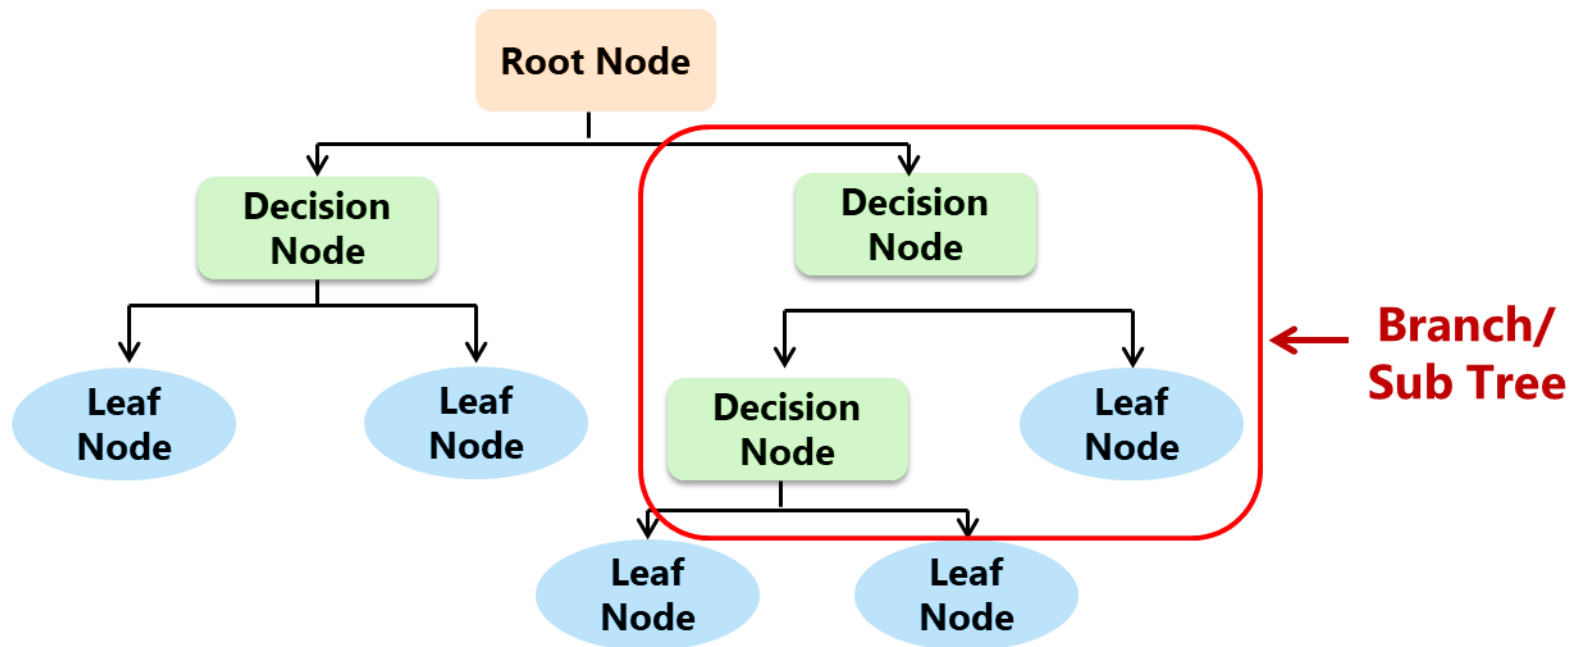
- 3 classes: A, B, C
- 1 instance per class
- 2 discrete features





# Decision trees: introduction

- **prediction model** (*non-parametric*); it does **not** learn a statistical model; it **learns an heuristic from known cases**
- **tree structure** with the following types of nodes:
  - **root node** is the top-most decision node and represents the entire set  $X$  that further gets divided according to the value of attributes
  - **internal decision node**  $f_m(x)$  based on an attribute  $m$ ; it produces as many branches as allowed by the splitting criterion adopted
  - **leaf node (or terminal node)**, with output value (class, prediction); it groups instances, for example, of a single class
- **decision**: for each new input instance  $x$ ,  $f_m(x)$  **guides the analysis of  $x$  from the root to a leaf** that provides the result (**class**)

# Decision trees: types of nodes



# Decision trees: learning (induction) of a decision tree



## Data set

2 classes:  $\circ$ ,  $\square$

2 dimensions:  $x_1, x_2$

## Decision tree

rectangular nodes  $\rightarrow$  decision nodes



circular nodes  $\rightarrow$  leaf nodes

1<sup>st</sup> split:  $x_1 > w_1$

2<sup>nd</sup> split:  $x_2 > w_2$  (orthogonal to 1<sup>st</sup>)

# Decision trees: classification

Given a point  $p = (w_1 + \delta, w_2 - \delta)$ ,  $\delta > 0$ , the induced tree classifies  $p$  as belonging to class  $\bigcirc$  (see **coloured path**): **it goes through the tree structure from the root node until reaching a leaf**





# Decision trees: interpretability


- decision tree  $\equiv$  set of rules **IF-THEN-ELSE**
- for example, the previous tree is equivalent to:

```
IF  $x_1 > w_1$  THEN  
    IF  $x_2 > w_2$  THEN  
        class =  $\square$   
    ELSE  
        class =  $\bigcirc$   
ELSE  
    class =  $\bigcirc$ 
```


# Univariate trees: introduction

the  $f_m(x)$  of each internal node is defined in terms of a single dimension (attribute)  $x_i$  (see *previous example*)

if  $x_i$  is **discrete**, for example:  $x_i$  is an attribute *colour* ...



if  $x_i$  is **numeric** (ordered), for example:  $x_i$  is an attribute *weight* ...



binary partition of space

$L = \{\text{people} \mid \text{weight} < 60\}$

$R = \{\text{people} \mid \text{weight} \geq 60\}$

# Univariate trees: tree learning algorithm

- tree induction  $\equiv$  tree learning
- algorithms of local search:
  - begin: root node with the whole set of instances
  - heuristic (recursive): to divide the set that arrives at a node  $m$  using the most discriminating attribute, that is, the one that generates more class homogeneous (i.e., pure) disjoint subsets
    - if the attribute is numeric, binary division; if the attribute is discrete, as many “children” (new nodes) as possible values
  - end: until get (labelled) pure nodes

# Univariate trees: impurity of a node

**objective:** evaluate the **degree of homogeneity** of the set of instances  $X$  that reaches a node  $m$

**example:** let  $C = \{\omega_1=\text{A}, \omega_2=\text{B}\}$  be the set of classes

- case 1:  $X^1 = \{(x_1, \text{A}), (x_2, \text{B}), (x_3, \text{B}), (x_4, \text{A}), (x_5, \text{A}), (x_6, \text{B})\}$
- case 2:  $X^2 = \{(x_1, \text{A}), (x_2, \text{B}), (x_3, \text{A}), (x_4, \text{A}), (x_5, \text{A}), (x_6, \text{A})\}$

**observations:**

- $X^1$  has 3 instances of each class (**heterogeneous**)
- $X^2$  has 5 instances of **A** and 1 of **B** (**almost homogeneous**)
- intuition:  $X^1$  is **more impure** than  $X^2$  or  $X^2$  is **more pure** than  $X^1$

# Univariate trees: an impurity measure

**impurity measure of a node  $m$** ; given:

- $N_m$ , the number of instances that reach  $m$
- $N_{i,m}$ , the number of instances in  $m$  that belong to class  $\omega_i$
- $P_{i,m} = N_{i,m}/N_m$ , the probability of class  $\omega_i$  in  $m$
- $I_m$ , **entropy**, a measure of impurity ( $c$  classes):

$$I_m = - \sum_{i=1}^c P_{i,m} \cdot \log_2 P_{i,m}$$

(Note:  $0 \cdot \log 0 = 0$ )

# Univariate trees: how to apply the measure

... it continues from previous example: suppose we can choose whether node  $m$  is reached by  $X^1$  or  $X^2$

- in both cases  $N_m = 6$
- in  $X^1$ :  $N_{1,m} = 3, N_{2,m} = 3 \Rightarrow p_{1,m} = 0.50, p_{2,m} = 0.50$
- in  $X^2$ :  $N_{1,m} = 5, N_{2,m} = 1 \Rightarrow p_{1,m} = 0.83, p_{2,m} = 0.17$

**evaluation of entropy (impurity measure):**

- in  $X^1$ ,  $I_m = - (0.5 \log_2 0.5 + 0.5 \log_2 0.5) = 1$  (+ *impure*)
- in  $X^2$ ,  $I_m = - (0.83 \log_2 0.83 + 0.17 \log_2 0.17) = 0.66$  (+ *pure*)

**if one could choose ... it would be better  $X^2$  to reach  $m$  because its entropy is lower (i.e. higher purity)**

# Univariate trees: generalization to 2 classes

**context:** impurity value at node  $m$  for 2 classes

- $N_m = N_{1,m} + N_{2,m} \Rightarrow$   
 $p_{1,m} + p_{2,m} = 1 \Rightarrow$   
 $p_{2,m} = 1 - p_{1,m}$
- $I_m$ , entropy for  $c = 2$  classes  
(let it be  $p \equiv p_{1,m}$ ):

$$I_m = -p \log_2(p) - (1 - p) \log_2(1 - p)$$

# Univariate tree: pure node (impurity 0)

a **node**  $m$  is **pure** if all instances reaching  $m$  are of the same class, that is, if  $N_{1,m} = 0$  or  $N_{2,m} = 0$ ; suppose  $N_{1,m} = 0$ , then

$$N_{1,m} = 0 \Rightarrow p_{1,m} = 0, \quad p_{2,m} = 1 \Rightarrow I_m = 0$$


a **node**  $m$  is **pure**  $\Leftrightarrow I_m = 0$  (impurity 0)

a **pure node** becomes a **leaf** that is assigned the class label of its instances



# Univariate trees: impurity of a split

- $K$  classes or categories
- $x$  is the attribute chosen to split an impure node



$$\sum_{j=1}^n N_{m,j} = N_m$$

$N_{m,j}^i$  instances of the  $N_{m,j}$  belong to  $\omega_i$  such that  $\sum_{i=1}^K N_{m,j}^i = N_{m,j}$   
 $p_{m,j}^i = \frac{N_{m,j}^i}{N_{m,j}}$  is the probability of  $\omega_i$  in the branch  $j$

**Impurity of a split:**

$$I_m^* = - \sum_{j=1}^n \frac{N_{m,j}}{N_m} \left( \sum_{i=1}^K p_{m,j}^i \log_2 p_{m,j}^i \right)$$

# Univariate trees: choice of the attribute; local optimization

## Local optimization:

in each node (**not pure**) in which it is intended to continue splitting the training set, the **most discriminating attribute** will be chosen, that is, the one that produces a split that is:

- the least impure  $\equiv$
- the most homogeneous  $\equiv$
- with the greatest class separation

# Univariate trees: general tree induction algorithm

```
generate_tree ( $T_{tra}$ ,  $\alpha$ ):  $m$   
  if entropy  $I_m(T_{tra}) \leq \alpha$ :  
     $m \leftarrow$  class label of the most represented class in  $T_{tra}$   
  else  
     $a =$  the_most_discriminating_attribute ( $T_{tra}$ )  
     $m \leftarrow a$   
    for each branch  $a_i$  of  $a$ :  
       $tra_i = \{x \in T_{tra} \mid x \text{ satisfies the condition of } a_i\}$   
       $h =$  generate_tree ( $tra_i$ ,  $\alpha$ )  
      add_child ( $m$ ,  $h$ )  
  return  $m$ 
```

# Univariate trees: function “the most discriminating attribute”

```
the_most_discriminating_attribute(tra): a
  mín_entr  $\leftarrow$  MAX
  for each attribute  $x_i$ :
    if  $x_i$  is discrete:
      divide tra into  $tra_1, \dots, tra_N$  based on the values of  $x_i$ 
       $e \leftarrow I_m^*(tra_1, tra_2, \dots, tra_N)$ 
      if  $e < mín\_entr$ :  $mín\_entr \leftarrow e$ ;  $a \leftarrow x_i$ 
    else //  $x_i$  is numeric
      for each possible binary division ( $tra_1, tra_2$ ):
         $e \leftarrow I_m^*(tra_1, tra_2)$ 
        if  $e < mín\_entr$ :  $mín\_entr \leftarrow e$ ;  $a \leftarrow x_i$ 
  return a
```

# Pruning: introduction

## Context

- it is common to find high overlapping or outliers
  - to generate a “pure” tree trying to learn from:
    - outliers
    - small sample size (small number of instances)
    - overlapping
- reduces the generalization of the model (overfitting)

## Alternative

- pre-pruning: do not split nodes with few instances while growing the tree
- post-pruning: after building the tree to its depth, remove unnecessary subtrees

# Pruning: pre-pruning

## Idea

- let  $m$  be a node and  $N_m$  the number of samples (instances)
- stop splitting the training set  
if  $m$  is pure ("*natural end*") or if  $N_m < \theta$  (pruning)
- label  $m$  with the majority class among the  $N_m$  instances

# Pruning: post-pruning

**Idea:** to generate 3 independent (and disjoint) sets for training, pruning, and test

- **training:** generate the complete tree, with all its pure leaves, using the training set
- **pruning** (it is part of the learning stage):
  - replace each subtree by a leaf node with the label of the majority instances covered by that subtree
  - prune the subtree if the “surrogate” leaf node does not worsen its performance with the pruning set
- **evaluation:** obtain a performance measure of the final (pruned) tree by classifying the test set

# Pruning: summary

Experiments show that:

- **post-pruning** produces **more accurate trees** (both classification and regression) than pre-pruning
- **pre-pruning** is **much faster** than post-pruning



# Decision trees: advantages

- **Comprehensive:** it is good for interpreting data in a highly visual way
- **Simplicity:** it is one of the simplest algorithms since it has no complex formulas or data structures

# Decision trees: disadvantages

- It is **computationally expensive**. At each node, the candidate split must be sorted before determining the best
- It is sometimes **unstable as small variations in data** might lead to the formation of a new tree